1. Problem and cost

Forms and contracts must be typed into systems of record by hand.

A human review costs 1 per field. A wrong value that is booked costs 100.

So auto-booking pays only if it is right more than 99 percent of the time

(break-even (E-1)/E = 99% at E = 100; 95% at E = 20; 99.8% at E = 500).

The dev set is mostly empty: 57.6% of the fields are null (absent).

Hallucinating a value for an absent field is the expensive mistake.

2. Pipeline (implemented, tested with mocks)

1. Full text, never truncated (CUAD and PDF text layer). Unreadable PDF -> null, no request.

2. Chunks of 45,000 characters with 3,000 overlap, every field in batches of at most 10.

3. Parallel-ready jobs, shared 429 stop, one retry, 60 s wait budget, reply cache.

4. Merge chunks per sample: conflicting values -> null, span pieces joined with ' | '.

5. Policy per sample: evidence must occur in the text; confidence must be finite in [0,1].

6. probly: categorical distribution over sample answers, entropy in nats (diagnostic).

7. Decision: tie or null winner -> null; otherwise min(agreement, lowest model confidence).

8. Threshold chosen only on dev_cal with the official scorer; all-abstain if no saving.

3. Measured result (6 dev_cal documents, K=1)

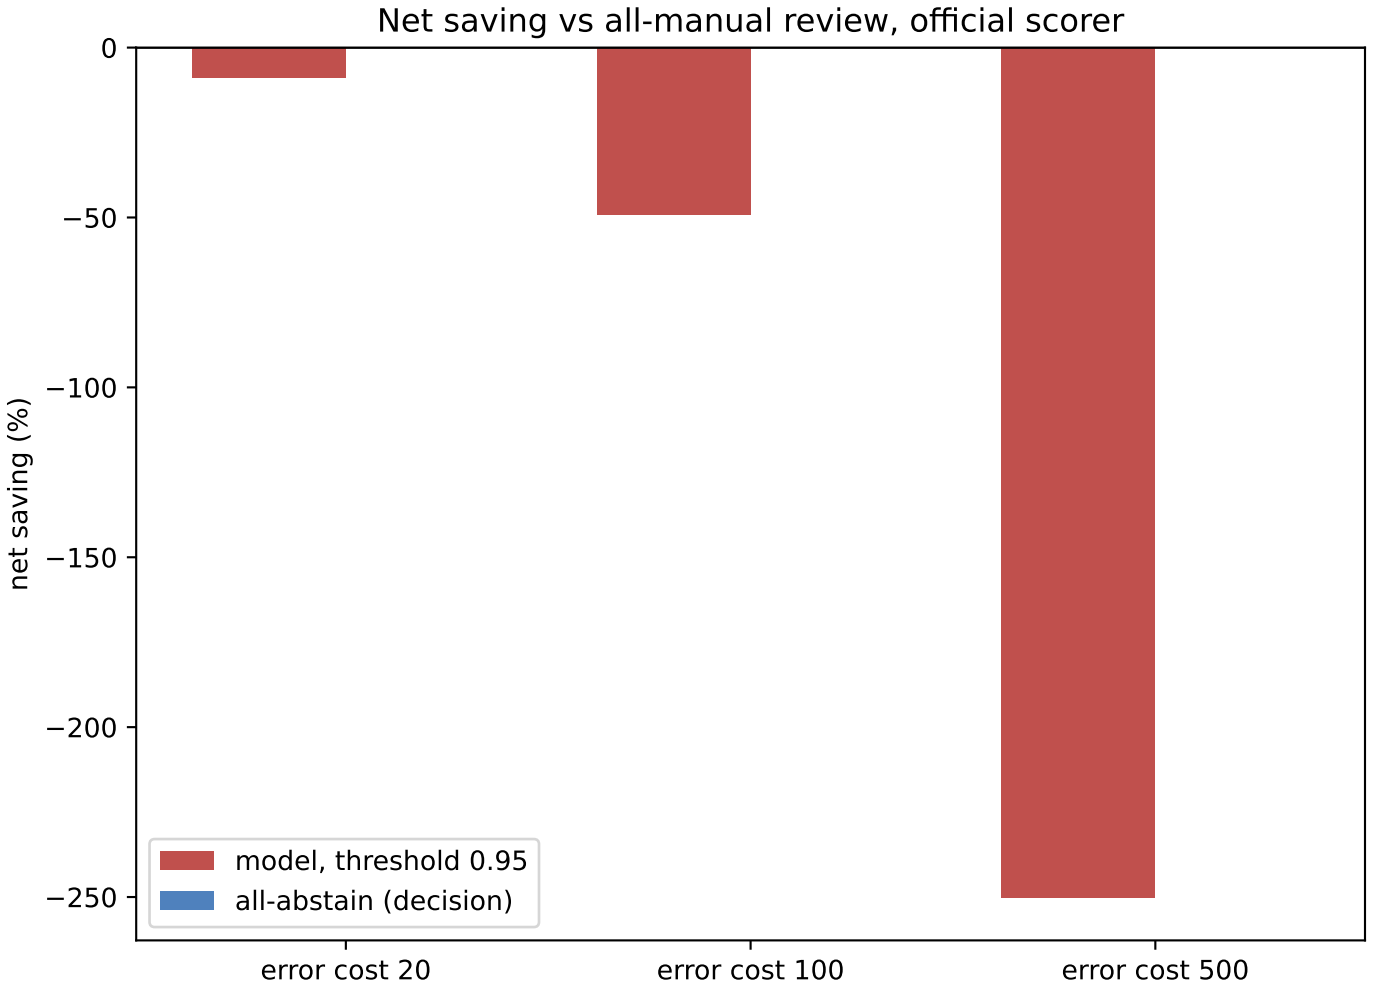

Functional test, not a quality proof: - 6 documents, 199 fields, 1 document completed fully. - 4 documents abstained after API 429 (team quota shared). - Model: precision 50% (2 accepted, 1 wrong), coverage 3.4%. - Decision: all-abstain. No positive saving was shown.

4. Accepted and abstained examples

Accepted, correct (confidence 0.95): Document Name = 'JOINT FILING AGREEMENT'.

Accepted, WRONG (confidence 0.95): Agreement Date = 'March 27, 2020'.

The gold keeps the final period: 'March 27, 2020.' The scorer's token F1 is 0.67 < 0.8.

Abstained by grounding (synthetic test): value 'Invented Value' with confidence 0.99.

It does not occur in the document, so the field is null with confidence 0.0.

Grounding shows that a value occurs in the text. It does not prove the right field:

a real text match in another section would still be accepted.

5. Limits, not tested, how to run

Not done: real K=3 runs (quota), Linux run (no Linux Python here), test_public (not used).

Confidence is an uncalibrated score. Entropy and agreement are diagnostics only.

Document-level abstention: one incomplete sample drops all fields of that document.

Strict spans cost recall; the scorer's token F1 penalises small punctuation changes.

Data flow: document text and prompts go to the external LLM endpoint (see SUBMISSION.md).

Run: bash run_hidden.sh <gold-free input> <fresh output> (decision: submission/decision.json)